

CS711008Z Algorithm Design and Analysis

Lecture 10. Algorithm design technique: Network flow and its applications¹

Dongbo Bu

Institute of Computing Technology
Chinese Academy of Sciences, Beijing, China

¹The slides are made based on Chapter 7 of Introduction to algorithms, Combinatorial optimization algorithm and complexity by C. H. Papadimitriou and K. Steiglitz. Some slides are excerpted from the presentation by K. Wayne with permission.

- MAXIMUMFLOW problem: FORD-FULKERSON algorithm, MAXFLOW-MINCUT theorem;
- A duality explanation of FORD-FULKERSON algorithm and MAXFLOW-MINCUT theorem;
- Scaling technique to improve FORD-FULKERSON algorithm;
- Solving the dual problem: Push-Relabel algorithm;
- Extensions of MAXIMUMFLOW problem: lower bound of capacity, multiple sources & multiple sinks, indirect graph;

A brief history of MINIMUMCUT problem I

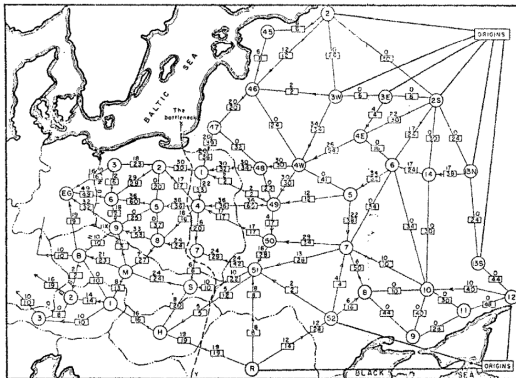


Figure: Soviet Railway network, 1955

A brief history of MINIMUMCUT problem II

- “.... From Harris and Ross [1955]: Schematic diagram of the railway network of the Western Soviet Union and Eastern European countries, with a maximum flow of value 163,000 tons from Russia to Eastern Europe, and a cut of capacity 163,000 tons indicated as The bottleneck.”
- A recently declassified U.S. Air Force report indicates that the original motivation of minimum-cut problem and Ford-Fulkerson algorithm is *to disrupt rail transportation the Soviet Union* [A. Shrijver, 2002].

A brief history of algorithms to MINIMUMCUT problem

Year	Developers	Time-complexity
1956	Ford and Fulkerson	$O(mC)$ and $O(m^2 \log C)$
1972	Edmonds and Karp	$O(m^2 n)$
1970	Dinitz	$O(n^2 m)$
1974	Karzanov	$O(n^3)$
1986	Sleator and Tarjan	$O(nm \log n)$
1988	Goldberg and Tarjan	$O(n^2 m \log(\frac{n^2}{m}))$
2012	Orlin	$O(nm)$

MAXIMUMFLOW problem

MAXIMUMFLOW problem

INPUT:


A directed graph $G = \langle V, E \rangle$. Each edge e has a capacity C_e .

Two special points: **source** s and **sink** t ;


OUTPUT:

For each edge $e = (u, v)$, to assign a flow $f(u, v)$ such that

$\sum_{u, (s, u) \in E} f(s, u)$ is maximized.



Intuition: to push as many commodity as possible from **source** s to **sink** t .



Definition (Flow)

$f : E \rightarrow \mathbb{R}^+$ is a **$s - t$ flow** if:


- ① (Capacity constraints): $0 \leq f(e) \leq C_e$ for all edge e ;
- ② (Conservation constraints): for any intermediate vertex $v \in V - \{s, t\}$, $f^{\text{in}}(v) = f^{\text{out}}(v)$, where
 $f^{\text{in}}(v) = \sum_{e \text{ into } v} f(e)$ and $f^{\text{out}}(v) = \sum_{e \text{ out of } v} f(e)$.
 (Intuition: input = output for any intermediate vertex.)

The **value of flow** f is defined as $V(f) = f^{\text{out}}(s)$.

Flow and Cut

Definition ($s - t$ cut)

An $s - t$ **cut** is a partition (A, B) of V such that $s \in A$ and $t \in B$. The **capacity of a cut** (A, B) is defined as $C(A, B) = \sum_{e \text{ from } A \text{ to } B} C(e)$.




$$C(A, B) = 3$$


Flow value lemma

Lemma

(Flow value lemma) Give a flow f . For **any** $s - t$ cut (A, B) , the flow across the cut is a constant $V(f)$. Formally, $V(f) = f^{out}(A) - f^{in}(A)$.



$$V(f) = 2 + 0 = 2$$
$$f^{out}(A) - f^{in}(A) = 2 + 1 - 1 = V(f)$$



Proof.

- We have: $0 = f^{out}(v) - f^{in}(v)$ for any node $v \neq s$ and $v \neq t$.
- Thus, we have:

$$\begin{aligned}
 V(f) &= f^{out}(s) - f^{in}(s) && // \text{Hint: } f^{in}(s) = 0; \\
 &= \sum_{v \in A} (f^{out}(v) - f^{in}(v)) \\
 &= \left(\sum_{e \text{ from } A \text{ to } B} f(e) + \sum_{e \text{ from } A \text{ to } A} f(e) \right) \\
 &\quad - \left(\sum_{e \text{ from } B \text{ to } A} f(e) + \sum_{e \text{ from } A \text{ to } A} f(e) \right) \\
 &= f^{out}(A) - f^{in}(A)
 \end{aligned}$$

FORD-FULKERSON algorithm [1956]

Lester Randolph Ford Jr. and Delbert Ray Fulkerson



Figure: Lester Randolph Ford Jr. and Delbert Ray Fulkerson

Trial 1: Dynamic programming technique

- Dynamic programming doesn't seem to work.
- In fact, there is no algorithm known for MAXIMUM FLOW problem that can really be viewed as belonging to the dynamic programming paradigm.
- We know that the MAXIMUMFLOW problem is in \mathbf{P} since it can be formulated as a linear program (See Lecture 8).
- However, the network structure has its own property to enable a more efficient algorithm, informally called **network simplex**, etc.

Trial 2: IMPROVEMENT strategy

Back to the general IMPROVEMENT strategy:

IMPROVEMENT(f)

```
1:  $\mathbf{x} = \mathbf{x}_0$ ; //starting from an initial solution;
2: while TRUE do
3:    $\mathbf{x} = \text{IMPROVE}(\mathbf{x})$ ; //move one step towards optimum;
4:   if STOPPING( $\mathbf{x}, f$ ) then
5:     break;
6:   end if
7: end while
8: return  $\mathbf{x}$ ;
```

Three key questions of iteration framework

Three key questions:

- ① How to construct an initial solution?
 - For MAXIMUMFLOW problem, a 0-flow can be obtained by setting $f(e) = 0$ for any e .
 - It is easy to verify that both CONSERVATION and CAPACITY constraints hold for the 0-flow.
- ② How to improve a solution?
- ③ When shall we stop?

A failure start: augmenting flow along a path in the original graph

- Let p be a simple $s - t$ path in the network G .
 - 1: Initialize $f(e) = 0$ for all e .
 - 2: **while** there is an $s - t$ path in graph G **do**
 - 3: **arbitrarily** choose an $s - t$ path p in G ;
 - 4: $f = \text{AUGMENT}(p, f)$;
 - 5: **end while**
 - 6: **return** f ;

Augmenting flow along a path


We define $bottleneck(p, f)$ as the minimum capacity of edges in path p .

AUGMENT(p, f) :


- 1: Let $b = bottleneck(p, f)$;
- 2: **for** each edge $e = (u, v) \in P$ **do**
- 3: **if** (u, v) is a forward edge **then**
- 4: increase $f(u, v)$ by b ;
- 5: **else**
- 6: decrease $f(u, v)$ by b ;
- 7: **end if**
- 8: **end for**

Why we fail?

- We start from 0-flow. In order to increase the value of f , we find a $s - t$ path, say $p = s \rightarrow u \rightarrow v$, to transmit more commodity.
- The flow on the three edges can be increased to 1 to meet both conservation and capacity constraints.
- However we cannot find a $s - t$ path in G to increase f further (left panel) although the maximum flow value is 2 (right panel).



$$V(f) = 1$$




$$V(f) = 2$$

Ford-Fulkerson algorithm: “undo” functionality

Key observation:



- When constructing a flow f , one might commit errors on some edges, i.e. the edges should not be used to transmit commodity. For example, the edge $u \rightarrow v$ should not be used.



- To improve the flow f , we should work out ways to **correct these errors**, i.e. “undo” the transmission assigned on the edges.

Implementing the "undo" functionality

- But how to implement the "undo" functionality?
- **Adding backward edges!**
- Suppose we add a **backwards** edge $v \rightarrow u$ into the original graph. Then we can correct the transmission via pushing back commodity from v to u .





Definition (Residual Graph)

Given a directed graph $G = \langle V, E \rangle$ with a flow f , we define **residual graph** $G_f = \langle V, E' \rangle$. For any edge $e = (u, v) \in E$, two edges are added into E' as follows:


- 1 (Forward edge (u, v) with leftover capacity):
If $f(e) < C(e)$, add edge $e = (u, v)$ with capacity $C(e) - f(e)$.
- 2 (Backward edge (v, u) with undo capacity):
If $f(e) > 0$, add edge $e' = (v, u)$ with capacity $C(e') = f(e)$.

Finding an $s - t$ path in G_f rather than G



Note: the path contains a backward edge (v, u)

Augmenting flow along the path: from f to f'



Note:

- By using the backward edge $v \rightarrow u$, the initial transmission from u to v is pushed back.
- More specifically, the first commodity transferred through flow f changes its path (from $s \rightarrow u \rightarrow v \rightarrow t$ to $s \rightarrow u \rightarrow t$), while the second one uses the path $s \rightarrow v \rightarrow t$.

FORD-FULKERSON algorithm

- Let p be a simple $s - t$ path in residual graph G_f , called **augmentation path**. We define $bottleneck(p, f)$ as the minimum capacity of edges in path p .

FORD-FULKERSON algorithm:


- 1: Initialize $f(e) = 0$ for all e .
- 2: **while** there is an $s - t$ path in residual graph G_f **do**
- 3: **arbitrarily** choose an $s - t$ path p in G_f ;
- 4: $f = \text{AUGMENT}(p, f)$;
- 5: **end while**
- 6: **return** f ;

Correctness and time-complexity analysis

Property 1: augmentation operation generates a new flow

Theorem

The operation $f' = \text{AUGMENT}(p, f)$ generates a new flow f' in G .



Proof.

- Checking **capacity constraints**: Consider two cases of edge $e = (u, v)$ in path p :
 - ① (u, v) is a forward edge arising from $(u, v) \in E$:
$$0 \leq f(e) \leq f'(e) = f(e) + \text{bottleneck}(p, f) \leq f(e) + (C(e) - f(e)) \leq C(e)$$
 - ② (u, v) is a backward edge arising from $(v, u) \in E$:
$$C(e) \geq f(e) \geq f'(e) = f(e) - \text{bottleneck}(p, f) \geq f(e) - f(e) = 0$$
- Checking **conservation constraints**:


On each node v , the change of the amount of flow entering v is the same as the change in the amount of flow exiting v .



Property 2: Monotonically increasing

Theorem

(Monotonically increasing) $V(f') > V(f)$




- Hint: $V(f') = V(f) + \text{bottleneck}(p, f) > V(f)$ since $\text{bottleneck}(p, f) > 0$.

Property 3: a trivial upper bound of flow

Theorem

$V(f)$ has an upper bound $C = \sum_{e \text{ out of } s} C(e)$.

(Intuition: the edges out of s are completely saturated with flow.)



Property 4: augmentation step

Theorem


Assume all edges have integer capacities. At every intermediate stage of the Ford-Fulkerson algorithm, both flow value $V(f)$ and residual capacities are integers. Thus, $bottleneck(p, f) \geq 1$, and there is at most C iterations of the while loop.

- Intuition: Under a reasonable assumption that all capacities are integers, we have $bottleneck(p, f) \geq 1$ at every stage; thus, $V(f') \geq V(f) + 1$.
- Time complexity: $O(mC)$. (Why? $O(C)$ iterations, and at each iteration, it takes $O(m + n)$ time to find an $s - t$ path in G_f using DFS or BFS technique.)


Property 5: A tighter upper bound

Theorem

(Tight upper bound) Given a flow f . For any $s - t$ cut (A, B) , we have $V(f) \leq C(A, B)$.



$$V(f) = 2 \leq C(A, B) = 3$$




Proof.

$$\begin{aligned}
 V(f) &= f^{out}(A) - f^{in}(A) && \text{(by flow value lemma)} \\
 &\leq f^{out}(A) && \text{(by } f^{in}(A) \geq 0) \\
 &= \sum_{e \in A \rightarrow B} f(e) \\
 &\leq \sum_{e \in A \rightarrow B} C(e) && \text{(by } f(e) \leq C(e)) \\
 &= C(A, B)
 \end{aligned}$$




Theorem

FORD-FULKERSON *ends up with a maximum flow f and a minimum cut (A, B) .*



Flow f



Residual graph G_f

Proof.

- FORD-FULKERSON algorithm ends when there is no $s - t$ path in the residual graph G_f .
- Let A be the set of nodes reachable from s in G_f , and $B = V - A$. (A, B) forms a $s - t$ cut. ($A \neq \phi, B \neq \phi$).
- Consider two types of edges $e = (u, v) \in E$ across cut (A, B) :
 - 1 $u \in A, v \in B$: we have $f(e) = C(e)$. Otherwise, A should be extended to include v since (u, v) is in G_f .
 - 2 $u \in B, v \in A$: we have $f(e) = 0$. Otherwise, A should be extended to include u since (v, u) is in G_f .
- Thus we have

$$\begin{aligned} V(f) &= f^{out}(A) - f^{in}(A) \\ &= f^{out}(A) \quad (\text{by } f^{in}(A) = 0) \\ &= \sum_{e \in A \rightarrow B} f(e) \\ &= \sum_{e \in A \rightarrow B} C(e) \quad (\text{by } f(e) = C(e)) \\ &= C(A, B) \end{aligned}$$


FORD-FULKERSON algorithm: bad example 1

The integer restriction is important


- In the analysis of FORD-FULKERSON algorithm, the integer restriction of capacities is important: the bottleneck edge leads to an increase of at least 1.
- The analysis doesn't hold if the capacities can be irrational.
- In fact, the flow might be increased by a smaller and smaller number and the iteration will be endless.
- Worse yet, this endless iteration might not converge to the maximum flow.

(See an example by Uri Zwick)


FORD-FULKERSON algorithm: bad example 2



A bad example of FORD-FULKERSON algorithm: Step 1




Flow $f : V(f) = 0$




An $s - t$ path in G_f

A bad example of FORD-FULKERSON algorithm: Step 2




Flow $f : V(f) = 1$



An $s - t$ path in G_f

A bad example of FORD-FULKERSON algorithm: Step 3



Flow $f : V(f) = 2$

Note:

- 1 After two iterations, the problem is similar to the original problem except for the capacities on (s, u) , (s, v) , (u, t) , (v, t) decrease by 1.
- 2 Thus, FORD-FULKERSON algorithm will end after $64+32$ iterations. (Why? *bottleneck* = 1 at all stages.)

FORD-FULKERSON algorithm: weakness

- FORD-FULKERSON algorithm doesn't specify how to choose an augmentation path, leading to some weaknesses:
 - A path with small bottleneck capacity is chosen as augmentation path;
 - We put flow on too many edges than necessary.
- The original max-flow paper also lists several heuristics for improvement.

Improvements of FORD-FULKERSON algorithm

- There are various implementations of the augmentation path selection:
 - ① Fat pipes:
 - To select the augmentation path with **the largest bottleneck capacity**;
 - Scaling technique: an efficient way to find an augmentation path with **large** improvement;
 - ② Short pipes:
 - Edmonds-Karp: to find **the shortest $s - t$ path** in *BFS tree*.
 - Dinitz' algorithm: to find a path in **layered network**, and perform **amortized** analysis;
 - Dinic's algorithm: running **DFS** to find a path in the **layered network constructed by running extended BFS**, and perform analysis using **blocking flow** technique;


Improvement 1: Scaling technique for speed-up

Scaling technique


- Question: can we choose a **large** augmentation path? The larger $bottleneck(p, f)$, the less iterations.
- An $s - t$ path p in G_f with the **largest** $bottleneck(p, f)$ can be found using binary search, or a slight change of Dijkstra's algorithm in $O(m + n \log n)$ time; however, it is still somewhat inefficient.
- Basic idea: we can relax the **“largest”** requirement to **“sufficiently large”**.
- Specifically, we can set up a lower bound Δ for $bottleneck(P, f)$: **simply removing the “small” edges**, i.e. the edges with capacities less than Δ from $G(f)$. This residual graph is called $G_f(\Delta)$.
- Δ will be scaled as iteration proceeds.

- Scaling FORD-FULKERSON algorithm:
 - 1: Initialize $f(e) = 0$ for all e .
 - 2: **Let** $\Delta = C$;
 - 3: **while** $\Delta \geq 1$ **do**
 - 4: **while** there is an $s - t$ path in $G_f(\Delta)$ **do**
 - 5: choose an $s - t$ path p ;
 - 6: $f = \text{AUGMENT}(p, f)$;
 - 7: **end while**
 - 8: $\Delta = \Delta/2$;
 - 9: **end while**
 - 10: **return** f ;
- Intuition: flow is augmented in a large step size whenever possible; otherwise, the step size is reduced. Step size is controlled via removing the “small” edges out of residual graph.
- Note: Δ turns to be 1 finally; thus no edge in residual graph will be neglected.

An example: Step 1




Flow $f : V(f) = 0$




No $s - t$ path in G_f

- Flow: 0 flow;
- Δ : $\Delta = 96$;
- $G_f(\Delta)$: the edges in light blue were removed since capacities are less than 96.
- $s - t$ path: cannot find. Thus Δ is scaled: $\Delta = \Delta/2 = 48$.

An example: Step 2




Flow $f : V(f) = 0$




An $s - t$ path in G_f

- Flow: 0 flow;
- Δ : $\Delta = 48$;
- $G_f(\Delta)$: the edges in light blue were removed since capacities are less than 48.
- $s - t$ path: a path $s - u - t$ appears. Perform augmentation operation.

An example: Step 3




Flow $f : V(f) = 64$




No $s - t$ path in G_f

- Flow: 64;
- Δ : $\Delta = 48$;
- $G_f(\Delta)$: the edges in light blue were removed since capacities are less than 48.
- $s - t$ path: no path found. Perform scaling: $\Delta = \Delta/2 = 24$.

An example: Step 4




Flow $f : V(f) = 64$




An $s - t$ path in G_f

- Flow: 64;
- Δ : $\Delta = 24$;
- $G_f(\Delta)$: the edges in light blue were removed since capacities are less than 24.
- $s - t$ path: find a path: $s - v - t$. Perform augmentation.

An example: Step 5



Flow $f : V(f) = 96$



No $s - t$ path in G_f

- Flow: 96. Maximum flow obtained.
- Δ : $\Delta = 24$;
- $G_f(\Delta)$: the edges in light blue were removed since capacities are less than 24.
- $s - t$ path: cannot find a $s - t$ path.

Theorem

(Outer while loop number) The while iteration number is at most $1 + \log_2 C$.

SCALING FORD-FULKERSON algorithm:

- 1: Initialize $f(e) = 0$ for all e .
- 2: **Let** $\Delta = C$;
- 3: **while** $\Delta \geq 1$ **do**
- 4: **while** there is an $s - t$ path in $G_f(\Delta)$ **do**
- 5: choose an $s - t$ path p ;
- 6: $f = \text{AUGMENT}(p, f)$;
- 7: **end while**
- 8: $\Delta = \Delta/2$;
- 9: **end while**
- 10: **return** f ;

Theorem

(Inner while loop number) In a scaling phase, the number of augmentations is at most $2m$.

SCALING FORD-FULKERSON algorithm:

- 1: Initialize $f(e) = 0$ for all e .
- 2: **Let** $\Delta = C$;
- 3: **while** $\Delta \geq 1$ **do**
- 4: **while** there is an $s - t$ path in $G_f(\Delta)$ **do**
- 5: choose an $s - t$ path p ;
- 6: $f = \text{AUGMENT}(p, f)$;
- 7: **end while**
- 8: $\Delta = \Delta/2$;
- 9: **end while**
- 10: **return** f ;

Analysis: inner while loop cont'd

Proof.

Notice that

- 1 Let f be the flow that a Δ -scaling phase ends up with, and f^* be the maximum flow. We have $V(f) \geq V(f^*) - m\Delta$. (Intuition: $V(f)$ is not too bad; the distance to maximum flow is small.)
- 2 In the subsequent $\frac{\Delta}{2}$ -scaling phase, each augmentation will increase $V(f)$ at least $\frac{\Delta}{2}$.

Thus, there are at most $2m$ augmentations in the $\frac{\Delta}{2}$ -scaling phase. \square

Time-complexity: $O(m^2 \log_2 C)$. (Reason: $O(\log_2 C)$ outer while loop, $O(m)$ inner loops, and each augmentation takes $O(m)$ time.)

But why $V(f) \geq V(f^*) - m\Delta$?

Proof.

- Let A be the set of nodes reachable from s in the residual graph $G_f(\Delta)$, and $B = V - A$. Thus (A, B) forms a cut ($A \neq \phi, B \neq \phi$).
- Consider two types of edges $e = (u, v) \in E$.
 - 1 $u \in A, v \in B$: we have $f(e) \geq C(e) - \Delta$. Otherwise, A should be extended to include v since (u, v) in $G_f(\Delta)$.
 - 2 $v \in A, u \in B$: we have $f(e) \leq \Delta$. Otherwise, A should be extended to include v since (u, v) in $G_f(\Delta)$, too.
- Thus we have:

$$\begin{aligned} V(f) &= \sum_{e \in A \rightarrow B} f(e) - \sum_{e \in B \rightarrow A} f(e) \\ &\geq \sum_{e \in A \rightarrow B} (C(e) - \Delta) - \sum_{e \in B \rightarrow A} \Delta \\ &\geq \sum_{e \in A \rightarrow B} C(e) - m\Delta \\ &= C(A, B) - m\Delta \\ &\geq V(f^*) - m\Delta \end{aligned}$$

Implementation 2: Edmonds-Karp algorithm using $O(m^2n)$ time

Edmonds-Karp algorithm [1972]



Figure: Jack Edmonds, and Richard Karp

Note: The algorithm was first published by Yefim Dinic in 1970 and independently published by Jack Edmonds and Richard Karp in 1972.

EDMONDS-KARP algorithm

EDMONDS-KARP algorithm:

- 1: Initialize $f(e) = 0$ for all e .
- 2: **while** there is a $s - t$ path in G_f **do**
- 3: choose **the shortest** $s - t$ path p in G_f using *BFS*;
- 4: $f = \text{AUGMENT}(p, f)$;
- 5: **end while**
- 6: **return** f ;

(a demo)

Theorem

Edmonds-Karp algorithm runs in $O(m^2n)$ time.

Proof.

- During the execution of Edmonds-Karp algorithm, an edge $e = (u, v)$ serves as **bottleneck** edge at most $\frac{n}{2}$ times
- Thus, the while loop will be executed at most $\frac{n}{2}m$ times since there are m edges in total
- It takes $O(m)$ time to find the shortest path using BFS, and augment flow using the path.





Lemma



An edge $e = (u, v)$ serves as a **bottleneck** edge at most $\frac{n}{2}$ times.

Proof.


- For a residual graph G_f , we first category all nodes into levels L_0, L_1, \dots , where $L_0 = \{s\}$, and L_i contains all nodes v such that the shortest path from s to v has i edges. We use $L(u)$ to denote the level number of node u .
- Consider the two consecutive occurrences of edge $e = (u, v)$ in G_f as bottleneck, say at step k and step k''' .
- We have $L(v) = L(u) + 1$ at step k , and after flow augmentation, the bottleneck edge $e = (u, v)$ will be reversed in G_f .
- At step k''' , $e = (u, v)$ becomes a **bottleneck** edge again.
- This means that $e' = (v, u)$ should be reversed first before step k''' , say at step k'' . We have $L''(u) = L''(v) + 1$. Thus $L''(u) = L''(v) + 1 \geq L'(v) + 1 \geq L(u) + 2$.
- For any node, its maximal level is at most n . Thus the lemma holds.











⋮





Implementation 3: Dinitz' algorithm and its variant Dinic's algorithm




Figure: Yefim Dinitz


The original Dinitz' algorithm

Motivations:


- The initial intention was just to accelerate FORD-FULKERSON algorithm by means of a smart data structure;
- Notice that finding an augmentation path takes $O(m)$ time and becomes a bottleneck of FORD-FULKERSON algorithm;
- It is valuable to save **all information** achieved at a BFS search for subsequent iterations;
- Specifically, the **BFS tree** is enriched to **layered network**:
 - BFS tree: includes **only the first edge found to a node v** ;
 - Layered network: keeping **all the edges residing on all the shortest $s - v$ path**;



Residual graph G_f



BFS tree



Layered network G_L

Dinic's algorithm: layered network + blocking flow

- Shimon Even and Alon Itai understood the paper by Y. Dinitz and that by A. Karzanov except for the **layered network maintenance** (removing the “dead-end” nodes). The gaps were spanned by using:
 - 1 **blocking flow** (first proposed by A. Karzanov) to prove that the levels of layered network increases from phase to phase;
 - 2 **DFS** to search an augmentation path.

DINIC's algorithm

```
1: Initialize  $f(e) = 0$  for all  $e$ .
2: while TRUE do
3:   Construct layered network  $G_L$  from residual graph  $G_f$ ;
4:   if  $\text{dist}(s, t) = \infty$  then
5:     break;
6:   end if
7:   find a blocking flow  $f'$  in  $G_L$  using DFS technique;
8:   augment flow  $f$  by  $f'$ ;
9: end while
10: return  $f$ ;
```


- Here, a **blocking flow** refers to a flow such that in the corresponding residual graph, there is no $s - t$ path.
- Intuition: after acquiring a layered network using $O(m)$ time, a blocking flow (containing several $s - t$ paths) is found for further augmentation. In contrast, EDMONDS-KARP algorithm augment only one path.

(See ppt for a demo)

- Constructing layered network: $O(m)$ (extended BFS)
- Finding blocking flow: $O(mn)$. The reason is:
 - ① it takes $O(n)$ time to find a $s - t$ path using DFS in a layered network;
 - ② at least one bottleneck edge in the path will be saturated;
 - ③ thus it needs at most m iterations to find a blocking flow;
- #WHILE = $O(n)$. (why? same argument to Edmonds-Karp analysis)
- Total: $O(mn^2)$

Understanding network-flow from the dual point of view


Duality explanation of MaxFlow-MinCut: Dual problem



DUAL: set variables for **edges** (Intuition: x_i denotes *flow* via edge i)

$$\begin{array}{llllll}
 \max & & & & & f \\
 \text{s.t.} & x_1 & +x_2 & & & -f = 0 \text{ vertex } s \\
 & & & -x_4 & -x_5 & +f = 0 \text{ vertex } t \\
 & -x_1 & & +x_3 & +x_4 & = 0 \text{ vertex } u \\
 & & -x_2 & -x_3 & & +x_5 = 0 \text{ vertex } v \\
 & x_1 & & & & \leq C_1 \\
 & & x_2 & & & \leq C_2 \\
 & & & x_3 & & \leq C_3 \\
 & & & & x_4 & \leq C_4 \\
 & & & & & x_5 \leq C_5 \\
 & x_1, & x_2, & x_3, & x_4, & x_5 \geq 0
 \end{array}$$

An equivalent version



$$\begin{array}{llllllll}
 \max & & & & & & f & \\
 s.t. & x_1 & +x_2 & & & & -f & \leq 0 \text{ vertex } s \\
 & & & & -x_4 & -x_5 & +f & \leq 0 \text{ vertex } t \\
 & -x_1 & & +x_3 & +x_4 & & & \leq 0 \text{ vertex } u \\
 & & -x_2 & -x_3 & & +x_5 & & \leq 0 \text{ vertex } v \\
 & x_1 & & & & & & \leq C_1 \\
 & & x_2 & & & & & \leq C_2 \\
 & & & x_3 & & & & \leq C_3 \\
 & & & & x_4 & & & \leq C_4 \\
 & & & & & x_5 & & \leq C_5 \\
 & x_1, & x_2, & x_3, & x_4, & x_5 & & \geq 0
 \end{array}$$

Note: the constraints (1), (2), (3), and (4) force

$-x_2 - x_3 + x_5 = 0$. So do other constraints.

Duality explanation of MaxFlow-MinCut: Primal problem

PRIMAL: set variables for **nodes**.

$$\begin{array}{rcccccccccccl}
 \min & & & & C_1 z_1 & +C_2 z_2 & +C_3 z_3 & +C_4 z_4 & +C_5 z_5 & & \\
 s.t. & y_s & & -y_u & +z_1 & & & & & & \geq 0 \\
 & y_s & & & & & +z_2 & & & & \geq 0 \\
 & & & y_u & -y_v & & & +z_3 & & & \geq 0 \\
 & & -y_t & +y_u & & & & & +z_4 & & \geq 0 \\
 & & -y_t & & +y_v & & & & & +z_5 & \geq 0 \\
 & -y_s & +y_t & & & & & & & & \geq 1 \\
 & y_s, & y_t, & y_u, & y_v, & z_1, & z_2, & z_3, & z_4, & z_5 & \geq 0
 \end{array}$$

Note:

- ❶ Since the constraints involves the difference among y_s, y_u, y_v and y_t , one of them can be fixed without effects. Here, we fix $y_s = 0$. Thus we have $y_t \geq 1$ (by the constraint $-y_s + y_t \geq 1$).
- ❷ Constraint (4) requires $z_4 \geq y_t - y_u$, and the objective is to minimize a function containing $C_4 z_4$, forcing $y_t = 1$.
- ❸ Constraint (1) requires $z_1 \geq y_u$, and the objective is to minimize a function containing $C_1 z_1$, forcing $z_1 = y_u$. So does constraint (2).

An equivalent version

PRIMAL: set variables for **nodes**.

$$\begin{array}{rcccccccc}
 \min & & & C_1 z_1 & +C_2 z_2 & +C_3 z_3 & +C_4 z_4 & +C_5 z_5 \\
 s.t. & -y_u & & +z_1 & & & & & = 0 \\
 & & -y_v & & +z_2 & & & & = 0 \\
 & y_u & -y_v & & & +z_3 & & & \geq 0 \\
 & y_u & & & & & +z_4 & & \geq 1 \\
 & & y_v & & & & & +z_5 & \geq 1 \\
 y_s & & & & & & & & = 0 \\
 & y_t & & & & & & & = 1 \\
 & & y_u, & y_v, & z_1, & z_2, & z_3, & z_4, & z_5 \geq 0
 \end{array}$$

Note: the coefficient matrix of constraints (3), (4) and (5) is totally uni-modular, implying the optimal solution is an integer solution.

An equivalent version

PRIMAL: set variables for **nodes**.

$$\begin{array}{rcll}
 \min & & C_1 z_1 & + C_2 z_2 & + C_3 z_3 & + C_4 z_4 & + C_5 z_5 & \\
 s.t. & -y_u & + z_1 & & & & & = 0 \\
 & & -y_v & + z_2 & & & & = 0 \\
 & y_u & -y_v & & + z_3 & & & \geq 0 \\
 & y_u & & & & + z_4 & & \geq 1 \\
 & & y_v & & & & + z_5 & \geq 1 \\
 y_s & & & & & & & = 0 \\
 & y_t & & & & & & = 1 \\
 & & y_u, & y_v, & z_1, & z_2, & z_3, & z_4, & z_5 & = 0/1
 \end{array}$$

MaxFlow-MinCut: strong duality

$$\begin{array}{rcll}
 \min & & C_1 z_1 & + C_2 z_2 & + C_3 z_3 & + C_4 z_4 & + C_5 z_5 & \\
 s.t. & -y_u & +z_1 & & & & & = 0 \\
 & -y_v & & +z_2 & & & & = 0 \\
 & y_u & -y_v & & +z_3 & & & \geq 0 \\
 & y_u & & & & +z_4 & & \geq 1 \\
 & & y_v & & & & +z_5 & \geq 1 \\
 y_s & & & & & & & = 0 \\
 y_t & & & & & & & = 1 \\
 & y_u, & y_v, & z_1, & z_2, & z_3, & z_4, & z_5 = 0/1
 \end{array}$$

Observations:

- Intuition of primal variables: if node i is in A , $y_i = 0$; and $y_i = 1$ otherwise.
- The primal problem is essentially to find a cut. (Note: $z_1 = 1$ iff $y_s = 0$ and $y_u = 1$, i.e., edge (s, u) is a cut edge.)
- By weak duality, we have $f \leq c$. This is exactly the MaximumFlow-MinimumCut theorem.

FORD-FULKERSON algorithm is essentially a primal-dual algorithm

Dual problem and DRP

- DUAL D: set variables for **edges**;

$$\begin{array}{rcll}
 \max & & f & \\
 s.t. & x_1 & +x_2 & -f \leq 0 \text{ vertex } s \\
 & & & -x_4 -x_5 +f \leq 0 \text{ vertex } t \\
 & -x_1 & & +x_3 +x_4 \leq 0 \text{ vertex } u \\
 & & -x_2 -x_3 & +x_5 \leq 0 \text{ vertex } v \\
 & x_1 & & \leq C_1 \\
 & & x_2 & \leq C_2 \\
 & & & x_3 \leq C_3 \\
 & & & & x_4 \leq C_4 \\
 & & & & & x_5 \leq C_5 \\
 & x_1, & x_2, & x_3, & x_4, & x_5 & \geq 0
 \end{array}$$

- Recall how to write *DRP* from *D*:
 - Replacing the right-hand side C_i with 0;
 - Adding constraints: $x_i \leq 1, f \leq 1$;
 - Keep only the tight constraints J . Here we category J into two sets, i.e. $J = J^S \cup J^E$, where $J^S = \{i | x_i = C_i \text{ in dual D}\}$, and $J^E = \{i | x_i = 0 \text{ in dual D}\}$. Intuitively, J^S denotes the saturated arcs, and J^E denotes the empty arcs.

DRP corresponds to flow-augmentation

- DRP:

$$\begin{array}{llllllll}
 \max & & & & & & f & \\
 s.t. & x_1 & +x_2 & & & & -f & = 0 \text{ vertex } s \\
 & & & & -x_4 & -x_5 & +f & = 0 \text{ vertex } t \\
 & -x_1 & & +x_3 & +x_4 & & & = 0 \text{ vertex } u \\
 & & -x_2 & -x_3 & & +x_5 & & = 0 \text{ vertex } v \\
 & & & x_i & & & & \leq 0 \quad i \in J^S \\
 & & & x_j & & & & \geq 0 \quad j \in J^E \\
 & x_1, & x_2, & x_3, & x_4, & x_5, & f & \leq 1
 \end{array}$$

- FORD-FULKERSON algorithm is essentially a primal_dual algorithm since for DRP,
 - $\omega_{OPT} = 0$ implies that optimal solution is found.
 - $\omega_{OPT} = 1$ implies a $s - t$ path (with unit flow) in residual graph G_f .
- Why G_f ? $x_i \leq 0, i \in J^S$ denotes a backward edge, $x_j \geq 0, j \in J^E$ denotes a forward edge, and for other edges, there is no restriction for x_i .

Push-relabel algorithm [A. V. Goldberg, R. E. Tarjan, 1986]

Push-relabel algorithm [A. V. Goldberg, R. E. Tarjan, 1986]

The push-relabel algorithm is one of the most efficient algorithms to compute a maximum flow. The general algorithm has $O(n^2m)$ time complexity, while the implementation with FIFO vertex selection rule has $O(n^3)$ running time, the highest active vertex selection rule provides $O(n^2\sqrt{m})$ complexity, and the implementation with Sleator's and Tarjan's dynamic tree data structure runs in $O(nm\log(n/m))$ time. In most cases it is more efficient than the Edmonds-Karp algorithm, which runs in $O(nm^2)$ time.

- Basic idea: Augmenting flow method maintains feasibility of the dual linear program, while push-relabel method maintains feasibility of “primal” problem.
 - 1 Ford-Fulkerson: set variables for edges. Update flow on edges until G_f has no $s - t$ path;
 - 2 Push-relabel: set variables for nodes. Update a pre-flow f , maintaining the property that G_f has no $s - t$ path, until f is a flow.

Push-relabel algorithm: pre-flow

Definition: f is a pre-flow if

- (Capacity condition): $f(e) \leq C(e)$;
- (Excess condition): for all node $v \neq s$,
$$E_f(v) = \sum_{e \text{ into } v} f(e) - \sum_{e \text{ out of } v} f(e) \geq 0;$$
- If no intermediate node has excess, then a pre-flow f becomes a flow.
- The only difficulty is: How to describe the “no $s - t$ path in G_f ” constraint? Using label.

Push-relabel algorithm: label

Definition: (Valid label) For each node $v \in V$ and a pre-flow f , we define its height $h(v)$ such that:

- $h(t) = 0$ and $h(s) = n$;
- For each edge (u, v) in the residual graph G_f , we have $h(v) \geq h(u) - 1$;

(Intuition: for an edge in G_f , its end cannot be too lower than its head.)

Theorem

There is no $s - t$ path in G_f if there exist valid labels.

Proof.

- Suppose there is a $s - t$ path in G_f .
- Notice that $s - t$ path contains at most $n - 1$ edges.
- Due to $h(s) = n$ and $h(u) \leq h(v) + 1$, the height of t should be great than 0. A contradiction with $h(t) = 0$.

Push-relabel algorithm

High-level idea:


- Initialization:
 - 1 preflow is set as: $f(s, u) = C(s, u)$, and $f(u, v) = 0$ for other edges;
 - 2 labels are set as: $h(s) = n$ and $h(v) = 0$ for others.
- Iteration: at each step, processing a node v with excess $E_f(v) > 0$ as follows:
 - 1 If there exists a lower neighbor: push excess to lower nodes;
 - 2 Otherwise: increase its height $h(v)$ while keeping labels valid.

Push-relabel algorithm cont'd


Push-relabel algorithm:

```
1:  $h(s) = n$ ;
2:  $h(v) = 0$ ; for any  $v \neq s$ ;
3:  $f(e) = C(e)$  for all  $e = (s, u)$ ;
4:  $f(e) = 0$ ; for other edges;
5: while there exists a node  $v$  with  $E_f(v) > 0$  do
6:   if there exists an edge  $(v, w) \in G_f$  s.t.  $h(v) > h(w)$ ; then
7:     //Push excess from  $v$  to  $w$ ;
8:     if  $(v, w)$  is a forward edge; then
9:        $e = (v, w)$ ;
10:       $bottleneck = \min\{E_f(v), C(e) - f(e)\}$ ;
11:       $f(e) + = bottleneck$ ;
12:    else
13:       $e = (w, v)$ ;
14:       $bottleneck = \min\{E_f(v), f(e)\}$ ;
15:       $f(e) - = bottleneck$ ;
16:    end if
17:  else
18:     $h(v) = h(v) + 1$ ; //Relabel node  $v$ ;
19:  end if
20: end while
```

A demo of push-relabel algo: initialization





A maximum-flow instance





Initial pre-flow f
and residual graph G_f



A demo of push-relabel algo: Step 1




A demo of push-relabel algo: Step 2





A demo of push-relabel algo: Step 3





A demo of push-relabel algo: Step 4



A demo of push-relabel algo: Step 5



A demo of push-relabel algo: Step 6



Theorem

Push-relabel algo keeps label valid. Therefore, it outputs a maximum flow when ends.

Proof.

(Induction on the number of push and relabel operations.)

- Push operation: the new f is still a pre-flow since the capacity condition still holds.

$Push(f, v, w)$ may add edge (w, v) into G_f . We have $h(w) < h(v)$. (pre-condition). Thus, the label is valid for the new G_f .

- Relabel operation: The pre-condition implies $h(v) \leq h(w)$ for any $(v, w) \in G_f$. $relabel(f, h, v)$ changes $h(v) = h(v) + 1$. Thus, the new $h(v) \leq h(w) + 1$.



Theorem

For any node v , $\#Relabel \leq 2n - 1$. Thus, the total label operation number is less than $2n^2$.

Proof.

- ① (Connectivity): For a node w with $E_f(w) > 0$, there should be a path from w to s in G_f .
(Intuition: node w obtain a positive $E_f(w)$ through a node v by $Push(f, v, w)$. This operation also causes edge (w, v) to be added into G_f . Thus, there should be a path from w to s .)
- ② (Upper bound of $h(v)$): $h(v) < 2n - 1$ since there is a path from v to s . The length of the path is less than $n - 1$, $h(s) = n$, and $h(v) \leq h(w) + 1$ for any edge (v, w) in G_f .



Two types of $Push$ operations:

- 1 Saturated push (s-push): if $Push(f, v, w)$ causes (v, w) removed from G_f .
- 2 Unsaturated push (uns-push): other pushes.

$$\#Push = \#s - push + \#uns - push.$$

Theorem

$$\#s - push \leq 2nm.$$

Proof.

Consider an edge $e = (v, w)$. We will show that during the execution of algo, (v, w) appears in G_f at most $2n$ times.

- (Removing): a saturated $Push(f, v, w)$ removes (v, w) from G_f . We have $h(v) = h(w) + 1$.
- (Adding): Before applying $Push(f, v, w)$ again, (v, w) should be added to G_f first. The only way to add (v, w) to G_f is $Push(f, w, v)$. The pre-condition of $Push(f, w, v)$ requires that $h(w) \geq h(v) + 1$, i.e., $h(w)$ should be increased at least 2 since the previous $Push(f, v, w)$ operation. And we have $h(w) \leq 2n - 1$.



Theorem

$$\#uns - push \leq 2n^2m.$$

Proof.

Define a measure $\Phi(f, h) = \sum_{v: E_f(v) > 0} h(v)$.

- (Increase and upper bound) $\Phi(f, h) < 4n^2m$:
 - ① Relabel: a relabel operation increase $\Phi(f, h)$ by 1. The total $O(2n^2)$ relabel operations increase $\Phi(f, h)$ at most $O(2n^2)$.
 - ② Saturized push: A saturated $Push(f, v, w)$ operation increases $\Phi(f, h)$ by $h(w)$ since w has excess now. $h(w) \leq 2n - 1$ implies an upper bound for each operation. The total $2nm$ saturated pushes increase $\Phi(f, h)$ by at most $4n^2m$.
- (Decrease) An unsaturated $Push(f, v, w)$ will reduce $\Phi(f, h)$ at least 1.

(Intuition: after unsaturated $Push(f, v, w)$, we have $E_f(v) = 0$, which reduce $h(v)$ from $\Phi(f, h)$; on the other side, w obtains excess from v , which will increase $\Phi(f, h)$ by $h(w)$. From $h(v) \leq h(w) + 1$, we have that $\Phi(f, h)$ reduces at least 1.)

Extension: an instance with multiple optimal solution

Multiple optimal solutions

- For some instances, there might be multiple optimal solutions, i.e. multiple flow with the same maximum flow value.
- In addition, these maximum flows might share some edges. In other words, these edges are “necessary” edges.
- Sometimes, we need to enumerate all these optimal solutions, or get a small sample of these optimal solutions.

(See ppt for a demo)